

Session 5 HW: Model Fine Tuning, Distillation and Pruning (Screenshots)

Project Summary:

- GPT-2 model is tuned using a mixture of synthetic optical module data and actual data from Cell Site Router
- Model is tuned for learning laser bias current, channel power of the optical module.
- The objective is to make a fine-tuned model from GPT-2, that is then distilled and pruned.
- The distilled pruned version of model is then quantized to run in 4bit mode in an embedded system with small foot print compute resources.
- Last step is to demonstrate an agent answering field technician's troubleshooting queries.

1. Model Fine Tuning

```
df.to_csv('synthetic_switch_telemetry.csv', index=False)
return df

df = generate_synthetic_switch_telemetry(200)
df.head()
```

```
[1]:
```

	timestamp	name	temp	trans-volt	channel_1_in_pwr	channel_1_out_pwr	channel_1_laser_bias_cur	channel_2_in_pwr	channel_2_out_pwr	channel_2_laser_bias
0	2025-03-08 00:44:21.737018	QSFP-1	42.472407	3.392609	-2.381257	-1.986390	8.536193	-1.889202	1.549579	5.8:
1	2025-03-08 00:39:21.737018	QSFP-2	77.042858	3.225242	2.415317	-1.328458	5.762695	0.251406	-2.852479	5.8:
2	2025-03-08 00:34:21.737018	QSFP-3	63.919637	3.248489	0.031514	-1.937937	7.881442	2.237675	-2.867259	5.1:
3	2025-03-08 00:29:21.737018	QSFP-4	55.919509	3.469566	1.958745	-2.467785	8.033575	1.393349	-1.058339	8.6:
4	2025-03-08 00:24:21.737018	QSFP-5	29.361118	3.381929	-1.079702	-2.276185	7.120653	1.839367	-0.068141	8.3:

```
with open('train.jsonl', 'w') as f:
    for _, row in df.iterrows():
        json.dump({"prompt": row_to_prompt(row), "response": classify_anomaly(row)}, f)
        f.write('\n')

print("✅ train.jsonl created.")
```

✅ train.jsonl created.

Session 5 HW: Model Fine Tuning, Distillation and Pruning (Screenshots)

```
print(f"PyTorch Version: {torch.__version__}")

try:
    teacher = AutoModelForCausalLM.from_pretrained("./fine_tuned_gpt2_telemetry", device_map=None)
    print("✅ Model loaded successfully to CPU.")

    teacher = teacher.to("cuda")
    print("✅ Model moved to GPU successfully.")
except Exception as e:
    print(f"❌ Error during model load/move: {e}")
```

CUDA Available: True
CUDA Device: Quadro M1000M
CUDA Version: 11.7
PyTorch Version: 2.0.0+cu117

C:\Miniconda3\envs\sentence-transformers\lib\site-packages\torchvision\datapoints_init_.py:12: UserWarning: The torchvision.datapoints and torchvision.transforms.v2 namespaces are still Beta. While we do not expect major breaking changes, some APIs may still change according to user feedback. Please submit any feedback you may have in this issue: <https://github.com/pytorch/vision/issues/6753>, and you can also check out <https://github.com/pytorch/vision/issues/7319> to learn more about the APIs that we suspect might involve future changes. You can silence this warning by calling torchvision.disable_beta_transforms_warning().
warnings.warn(_BETA_TRANSFORMS_WARNING)
C:\Miniconda3\envs\sentence-transformers\lib\site-packages\torchvision\transforms\v2_init_.py:54: UserWarning: The torchvision.datapoints and torchvision.transforms.v2 namespaces are still Beta. While we do not expect major breaking changes, some APIs may still change according to user feedback. Please submit any feedback you may have in this issue: <https://github.com/pytorch/vision/issues/6753>, and you can also check out <https://github.com/pytorch/vision/issues/7319> to learn more about the APIs that we suspect might involve future changes. You can silence this warning by calling torchvision.disable_beta_transforms_warning().
warnings.warn(_BETA_TRANSFORMS_WARNING)
✅ Model loaded successfully to CPU.
✅ Model moved to GPU successfully.

```
import os
import torch
from transformers import (
    AutoModelForCausalLM,
    AutoTokenizer,
    TrainingArguments,
    Trainer,
    DataCollatorForLanguageModeling
)
from datasets import load_dataset

# --- Debug GPU Information ---
print(f"CUDA Available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"CUDA Device: {torch.cuda.get_device_name(0)}")
    print(f"CUDA Version: {torch.version.cuda}")
    print(f"PyTorch Version: {torch.__version__}")
    os.environ['CUDA_LAUNCH_BLOCKING'] = '1' # Force clearer error reporting from CUDA

# --- Load Dataset ---
dataset = load_dataset('json', data_files={'train': 'train.jsonl'})
train_test_split = dataset['train'].train_test_split(test_size=0.2)
train_dataset = train_test_split['train']
eval_dataset = train_test_split['test']

# --- Load Tokenizer ---
model_name = 'gpt2'
tokenizer = AutoTokenizer.from_pretrained(model_name)
```

Session 5 HW: Model Fine Tuning, Distillation and Pruning (Screenshots)

2. Checking the fine tuned model for corruption etc.

▼ Load Fine Tuned Model for Distillation

```
[1]: import torch
from transformers import AutoModelForCausalLM, AutoTokenizer

print(f"CUDA Available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"CUDA Device: {torch.cuda.get_device_name(0)}")
    print(f"CUDA Version: {torch.version.cuda}")
    print(f"PyTorch Version: {torch.__version__}")

try:
    teacher = AutoModelForCausalLM.from_pretrained("./fine_tuned_gpt2_telemetry", device_map=None)
    print("✅ Model loaded successfully to CPU.")

    teacher = teacher.to("cuda")
    print("✅ Model moved to GPU successfully.")
except Exception as e:
    print(f"❌ Error during model load/move: {e}")
```

CUDA Available: True
CUDA Device: Quadro M1000M
CUDA Version: 11.7
PyTorch Version: 2.0.0+cu117



C:\Miniconda3\envs\sentence-transformers\lib\site-packages\torchvision\datapoints_init_.py:12: UserWarning: The torchvision.datapoints and torchvision.transforms.v2 namespaces are still Beta. While we do not expect major breaking changes, some APIs may still change according to user feedback. Please submit any feedback you may have in this issue: <https://github.com/pytorch/vision/issues/6753>, and you can also check out <https://github.com/pytorch/vision/issues/7319> to learn more about the APIs that we suspect might involve future changes. You can silence this warning by calling `torchvision.disable_beta_transforms_warning()`.

3. Model Distillation

```
# --- Save Distilled Student Model ---
student.save_pretrained("distilled_gpt2_telemetry")
tokenizer.save_pretrained("distilled_gpt2_telemetry")

print("✅ CPU-only Distillation complete - Student saved to 'distilled_gpt2_telemetry'")
```

WARNING:tensorflow:From C:\Miniconda3\envs\sentence-transformers\lib\site-packages\tf_keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

Map: 100%  160/160 [00:00<00:00, 275.92 examples/s]
Map: 100%  40/40 [00:00<00:00, 204.61 examples/s]

C:\Miniconda3\envs\sentence-transformers\lib\site-packages\transformers\training_args.py:1575: FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.46 of Transformers. Use `eval_strategy` instead
warnings.warn(
C:\Miniconda3\envs\sentence-transformers\lib\site-packages\transformers\training_args.py:1590: FutureWarning: using `no_cuda` is deprecated and will be removed in version 5.0 of Transformers. Use `use_cpu` instead
warnings.warn(
C:\Users\shaji.nathan\AppData\Local\Temp\ipykernel_8464\3197698090.py:121: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `DistillationTrainer.__init__`. Use `processing_class` instead.
trainer = DistillationTrainer(
[480/480 4:28:27, Epoch 3/3]

Epoch	Training Loss	Validation Loss
1	33.666700	26.494558
2	26.089500	19.545992
3	23.715200	19.037495

✅ CPU-only Distillation complete - Student saved to 'distilled_gpt2_telemetry'

Model is saved in the distilled_gpt2_telemetry directory

Session 5 HW: Model Fine Tuning, Distillation and Pruning (Screenshots)

4. Do a quick check for corruption etc.

Load Distilled Model into CPU and GPU to check for corruption

```
[5]: import torch
from transformers import AutoModelForCausalLM, AutoTokenizer

print(f"CUDA Available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"CUDA Device: {torch.cuda.get_device_name(0)}")
    print(f"CUDA Version: {torch.version.cuda}")
    print(f"PyTorch Version: {torch.__version__}")

try:
    teacher = AutoModelForCausalLM.from_pretrained("./distilled_gpt2_telemetry", device_map=None)
    print("✅ Model loaded successfully to CPU.")

    teacher = teacher.to("cuda")
    print("✅ Model moved to GPU successfully.")
except Exception as e:
    print(f"❌ Error during model load/move: {e}")

CUDA Available: True
CUDA Device: Quadro M1000M
CUDA Version: 11.7
PyTorch Version: 2.0.0+cu117
✅ Model loaded successfully to CPU.
✅ Model moved to GPU successfully.
```

5. Prune the distilled Model

▼ Model Pruning

```
]: from torch.nn.utils import prune

model = AutoModelForCausalLM.from_pretrained("./distilled_gpt2_telemetry")

for name, module in model.named_modules():
    if isinstance(module, torch.nn.Linear):
        prune.l1_unstructured(module, name='weight', amount=0.3)

model.save_pretrained("distilled_pruned_gpt2_telemetry")
print("✅ Model pruned and saved as 'distilled_pruned_gpt2_telemetry'.")

✅ Model pruned and saved as 'distilled_pruned_gpt2_telemetry'.
```

Session 5 HW: Model Fine Tuning, Distillation and Pruning (Screenshots)

6. Check the distilled/pruned model by loading into the CPU and GPU

▼ Check Distilled and Pruned Model for corruption

```
8]: import torch
from transformers import AutoModelForCausalLM, AutoTokenizer

print(f"CUDA Available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"CUDA Device: {torch.cuda.get_device_name(0)}")
    print(f"CUDA Version: {torch.version.cuda}")
    print(f"PyTorch Version: {torch.__version__}")

try:
    teacher = AutoModelForCausalLM.from_pretrained("./distilled_pruned_gpt2_telemetry", device_map=None)
    print("✅ Model loaded successfully to CPU.")

    teacher = teacher.to("cuda")
    print("✅ Model moved to GPU successfully.")
except Exception as e:
    print(f"❌ Error during model load/move: {e}")

CUDA Available: True
CUDA Device: Quadro M1000M
CUDA Version: 11.7
PyTorch Version: 2.0.0+cu117

Some weights of the model checkpoint at ./distilled_pruned_gpt2_telemetry were not used when initializing GPT2LMHeadModel: ['lm_head.weight_mask']
- This IS expected if you are initializing GPT2LMHeadModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing GPT2LMHeadModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).
✅ Model loaded successfully to CPU.
✅ Model moved to GPU successfully.
```

Directory of the Models Created

Search finetune				
Name	Date modified	Type	Size	
.ipynb_checkpoints	3/8/2025 10:34 AM	File folder		
distilled_gpt2_telemetry	3/8/2025 3:07 PM	File folder		
distilled_pruned_gpt2_telemetry	3/8/2025 3:12 PM	File folder		
distilled_results	3/8/2025 3:00 PM	File folder		
fine_tuned_gpt2_telemetry	3/8/2025 1:00 AM	File folder		
logs	3/8/2025 10:38 AM	File folder		
quantized_pruned_gpt2_telemetry	3/8/2025 3:11 PM	File folder		
results	3/8/2025 1:00 AM	File folder		

Session 5 HW: Model Fine Tuning, Distillation and Pruning (Screenshots)

7. Quantize the distilled and pruned Model for edge device with small CPU/Memory footprint

```
# Step 8: Save the Quantized Model Locally
save_directory = "../quantized_distilled_pruned_gpt2_telemetry"
print(f"💎 Saving the quantized model to {save_directory}...")
model.save_quantized(save_directory)
tokenizer.save_pretrained(save_directory) # Save tokenizer
print(f"✅ Quantized model and tokenizer saved successfully in {save_directory}!")

💎 Saving the quantized model to ../quantized_distilled_pruned_gpt2_telemetry...
✅ Quantized model and tokenizer saved successfully in ../quantized_distilled_pruned_gpt2_telemetry!

# Step 9: Verify Saved Files
import os
print("\n📁 Saved files:")
print(os.listdir(save_directory))

📁 Saved files:
['config.json', 'gptq_model-4bit-128g.safetensors', 'merges.txt', 'quantize_config.json', 'special_tokens_map.json', 'tokenizer.json', 'tokenizer_config.json', 'vocab.json']

# Step 10: Reload the Saved Quantized Model
print("\n💎 Reloading the quantized model for verification...")
reloaded_model = AutoGPTQForCausalLM.from_quantized(save_directory)
reloaded_model.to(device)
print(f"✅ Quantized model reloaded successfully!")

WARNING - Exllamav2 kernel is not installed, reset disable_exllamav2 to True. This may because you installed auto_gptq using a pre-build wheel on Windows, in which exllama_kernels are not compiled. To use exllama_kernels to further speedup inference, you can re-install auto_gptq from source.
WARNING - CUDA kernels for auto_gptq are not installed, this will result in very slow inference speed. This may because:
1. You disabled CUDA extensions compilation by setting BUILD_CUDA_EXT=0 when install auto_gptq from source.
2. You are using pytorch without CUDA support.
3. CUDA and nvcc are not installed in your device.
WARNING - ignoring unknown parameter in quantize_config.json: quant_method.
INFO - The layer lm_head is not quantized.
```

8. AiOPS agent in action, using the quantized Model

```
NetAdmin, type your network issue. Type 'exit' to quit.
👉 Describe the problem (or type 'exit' to quit): Optical Module not working
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.

🤖 AiOPS Assistant Response:
[AiOPS Assistant] NetAdmin request: Optical Module not working
[AiOPS Response]: Anomalous - Channel 1 Input Power Over-Coupling Timestamp 0.21 dBm (SFP)
response: Normal, QPI+2 Output Voltage 2.19V
Response Time 3.62dBm (-1.74 dB m²), Max B

-----
👉 Describe the problem (or type 'exit' to quit): exit

🔴 Exiting AiOPS Assistant. Have a great day, NetAdmin! 🚀
```